

~~DT
10/Nov/2020~~

Assignment - 13

Name - Avneet Kumar Agarwal
Branch - CSE-I
Roll No - 1602040031

Chapter - 14

Q 14.1 omission of any of the 3 requirements of Defn 14.3 (termination, agreement, non-triviality) for the consensus problem allows a very simple form. Show that by presenting the 3 simple forms.

- if only termination & agreement are required:
Each process immediately decides on 0. (This algorithm is trivially correct no decision is possible).
- if only termination & non-triviality are required:
Each process immediately decides on its input. (Agreement is not satisfied because different decisions can be taken.)
- if only agreement & non-triviality are required:
Each process ~~should decide~~ should shout its input, then wait until $\lceil \frac{N+1}{2} \rceil$ messages with the same value have been received, & decides on this value. (This algorithm doesn't decide in every correct process if insufficiently many processes shout the same value.)

Ques. Give deterministic, crash robust consensus protocols for each for the following restrictions on the input values.

- (1) It is given that the parity of the input is even. (i.e., there are an even no. of processes with input 1) in each initial configuration.
- (2) There are 2 (known) processes ~~or~~ or n_0, n_1 , each initial config satisfies $n_0 = n_1$.
- (3) In each initial configuration, there are at least $\lceil \frac{n}{2} + 1 \rceil$ processes with the same input.

Ans (1) The even parity implies that the entire array can be reconstructed if $n-1$ inputs are known.
Since every process shouts its input & waits for the receipt of $n-1$ values (this includes the process' own input). The n^{th} value is chosen so as to make the overall parity even. The process decides on the most frequent input (or if there is a draw).

(2) So ≥ 2 processes, n_0, n_1 shout their inputs.

A process decides on the first value received.

(3) So ≥ 1 Each process shouts its input, awaits the receipt of $n-1$ values, & decides on the most frequently received value.

After this, a process sends every value.

Q4.3 Show that there is no $t = \text{initially dead - notion}$ election algorithm for $t \geq N/2$.

2 sets of processes $S \& T$ can be formed such that the processes in each group decide independently of the other group. If one group can reach a ~~definition~~ decision in which there is a leader, an execution can be constructed in which 2 leaders are elected. If one group can reach a decision in which no leader is elected, an execution can be constructed in which all processes are correct, but no leader is elected.

Q4.4 Show that no algorithm for ϵ -approximate agreement can tolerate $t \geq N/2 + \lfloor \epsilon \rfloor$.

choose disjoint subsets $S \& T$ of $N-t$ processes & give all processes in S input 1 & the processes in T input $1+2\epsilon$. The processes in S can reach a decision on their own, because all inputs in this group is 1, so are the outputs. Indeed, the same sequence of steps is applicable if all process have input = 1, in which case 1. is the only allowed output. Similarly, the processes in T decide on output $1+2\epsilon$ on their own. The decisions of $S \& T$ can be taken in the same run, contradicting the agreements requirement.

M.8 Give a bijection from the set
 $\{(s, t) : N-t \leq s \leq N \text{ & } 1 \leq t \leq s\}$
 to the integers in the range $[1, \dots, k]$.

Ans $f(s, t) = \frac{1}{2}(N-t+s-1)(s-(N-t)) + t,$

M.9 Adapt the proof of Theorem M.15 for the case that G_T consists of k -connected components.

M.15 \rightarrow there exists no non-trivial t -coash-robust deletion algorithm for a disconnected task T ,
 → all the components are through G_T . After determining that the deletion vector belongs to a_i^* , decide on the parity of i .

e.g. consider the problem of (k, l) -election, which generalizes the usual election problem. The problem requires that all correct processes decide on either 0 ("elected" or 1 ("elected"), & the no. of processes that decide 1 $\leq k-l$ (inclusive).

What are the uses of (k, l) -election?

(1) What are the uses of (k, l) -election?

(2) Demonstrate that no deterministic t -coash robust algo. for (k, k) -election exists. ($\text{NP} \subseteq \text{CN}$)

(3) Give a deterministic t -coash robust algorithm for $(k, k+2)$ -election.

Ans (1) sometimes a leader is elected to coordinate a centralized algorithm. Execution of a small no. of steps of this algo. (e.g. in situations where election is not achievable) may still be considerably more efficient than execution by all processes.

(a) The decision vectors at (k_1, k_2) -election all have exactly k 1's, which implies that the table is disconnected (each node in cut is isolated). The problem is non-trivial because a crashed nodes can't decide 1; therefore, in each election the algorithm must decide on a vector where the 1's are located in correct positions. Consequently, the 14.15 applies.

(b) Every node should its identity, awaits the receipt of $N-t$ identities (including its own), & computes its rank γ in the set of $N-t$ identities. If $\gamma \leq k$ then it decides 1, otherwise it decides 0.

Q.M.9 (i) Does the convergence requirement imply that the expected no. of steps is bounded?

(ii) Is the expected no. of steps bounded in all algorithms of this section?

Ans (i) NO, it doesn't. The probability distribution on K defined by $P_0[K \geq k] = 1/k$ satisfies $\lim_{K \rightarrow \infty} P_0[K \geq k] = 0$

$E[K]$ is unbounded,

$= \infty$, but yet $E[K]$ is unbounded

(ii) In all algorithms the probability distribution on the no. of rounds K is geometric, i.e., there is a constant $c < 1$ such that $P_0[K \geq k] \leq c \cdot P_0[K \geq k]$ which implies that $E[K] \leq (1 - c)^{-1}$.

Q.10 Show that if all correct processes start round k of the crash-restart consensus algorithm (Alg M₃) then all correct processes will also finish round k if all correct processes start round k, at least $\frac{N-t}{2}$ processes finish in that round, allowing every correct process to receive N-t messages & finish the round.

M.3 (1) Prove that if more than $(N-t)/2$ processes start the crash-restart consensus algorithm (Alg M₃)

with input v, then a decision for v is taken in 3 rounds -

(2) Prove that if more than $(N-t)/2$ processes start

the algo, with $v \neq v'$, then a decision is possible.

(3) Is a decision for v possible if exactly

$(N-t)/2$ processes start the algorithm with $v \neq v'$

(4) What are the bivalent IP configurations at the end of the algorithm?

(5) In the first round, less than $\frac{(N-t)}{2}$ processes

vote for a value other than v, and implies that every (correct) process received a majority of votes for v. Hence, in round 2, only v-votes are submitted, so in round 3, all processes vote for v, to decide on v.

(6) Assuming the processes with $v \neq v'$ don't crash, more than $\frac{N-t}{2}$ processes vote for v, in 1st round.

It is possible that enough (correct) process receives all v-votes, & less than $\frac{N-t}{2}$ votes for the other value.

Then in round 2, only V-votes are submitted, implying that in round 3 all (correct) processes witness for V, & decide on V.

(3) Assuming the protocol with RIP V doesn't crash, exactly $\frac{N-t}{2}$ processes vote for V in the 1st round. In the most favourable (for V) case, an honest process receives $\frac{N-t}{2}$ V-votes & $\frac{N-t}{2}$ votes for the other value as well.

As seen from algo 14.3, the value '1' is adopted for the next round in case of a tie. Hence, a decision for V is possible if $t \geq 1$.

(4) An initial configuration is bivalent iff the no. of 1's, (#1), satisfies:

$$\left(\frac{N-t}{2} \right) \leq \#1 \leq \frac{N+t}{2}$$

Q14.12 (1) Prove that, if more than $(N+t)/2$ processes start algo 14.5, with RIP V, a V-decision is eventually taken.

(2) Prove that, if more than $(N+t)/2$ processes start algo 14.5, with RIP V, & $t < N/2$, a V-decision is taken in 2 rounds.

Q14.13 In the 1st round, more than $\frac{(N-t)}{2}$ correct processes vote for V, \rightarrow every process = $\frac{(N-t)}{2}$ correct processes choose V again in the 1st round, & hence chooses V. As in proof of lemma 14.3, it is shown that all correct processes choose V again in an later rounds, implying that a V-decision will be taken.

(2) As in point (1), in the 1st round all correct processes choose V_0 in 1st round, & vote for it in 2nd round. Hence, in the 2nd round each correct process selects at least $N-2t$ V-votes, i.e., $N-2t \leq N/2$, i.e., $N-2t \geq \frac{N+t}{2}$, so, it follows that in 2nd round every correct process decides on V .

M.13 Proved that no asynchronous t-Byzantine robust ~~also~~ broadcast algorithm exists for $t \geq N/3$.

AS Assume that such a protocol exists. Partition the process in 3 groups: S , $T \cup U$; each of size $\leq t$, with the general in S .
 \rightarrow Let δ_0 be the initial config' where the general has $\text{val} = V$.

\rightarrow Because all processes in U can be faulty, $\Rightarrow S \rightarrow SUT\delta_0$, whereas all processes in $S \cup T$ have decided on 0. Similarly, $T \rightarrow SUV\delta_0$, where in δ_0 all processes in SUT have decided 0.

\rightarrow Now, assume the processes in S are all faulty. & δ_0 is the initial configuration. First the processes in S cooperate with the processes in T such ~~as~~ away that all processes in SUT decide on 0, then the processes in T restore their state as in δ_0 & cooperate with the processes in U in such a way that all the processes in SUV decide on 1. Now the correct processes in $S \cup T$ have decided differently, which contradicts the statement/agreement.

Prove that during the execution of algo 14.6 at most $N(3N+1)$ messages are sent by correct processes.

At most N initial messages are sent by a correct process, namely by the general (if it is correct). Each ^{correct} process sends at most one echo, counting for N messages in each ~~for~~ correct process. Each correct process sends at most 2 ready messages, counting for $2N$ messages in each correct process. The no. of correct processes can be as high as N , resulting in the $N(3N+1)$ bound.

(Adapting the algorithm so as to suppress the sending of ready messages by correct processes improves the complexity to $N(6N+1)$.)

Chapter 15

Q15.1 How many messages are exchanged in the broadcast (N,t) protocol?

Ans Let $t \leq n = M_t(N)$,

$$M_t(N) \geq N-1.$$

$$M_t(N) = (N-1) + (N-1) \cdot M_{t-1}(N-1)$$

→ The 1st term is called initial transmission by the general, 2nd term represents the $N-1$ recursive calls, The recursion is solved by

$$M_t = \sum_{j=1}^t M_{t-j}(N-j)$$

Q15.2 What is the highest no. of messages sent by correct processes in Alg 15.2 if t executions that decide on 2? Answer both for the case where the general is correct & the case where the general is faulty.

Ans If decision = 0, while general is correct
the general's input is 0.

→ No correct process p initiates or supports any correct process, hence the messages sent by p are supporting a faulty process q; so p shouts at most t times; which means sending $(N-1)t$ messages.

→ If the general is faulty, at most $L-1 = t$ correct processes can initiate without enacting a 1-decision, correct process p shouts at most 1 $\langle bm, 1 \rangle$ message, at most $t \langle bm, 1 \rangle$ messages for faulty q, & at most $t \langle bm, 0 \rangle$ messages for ~~faulty~~ correct p; this means sending at most $(N-1)/2t + 1$ messages.

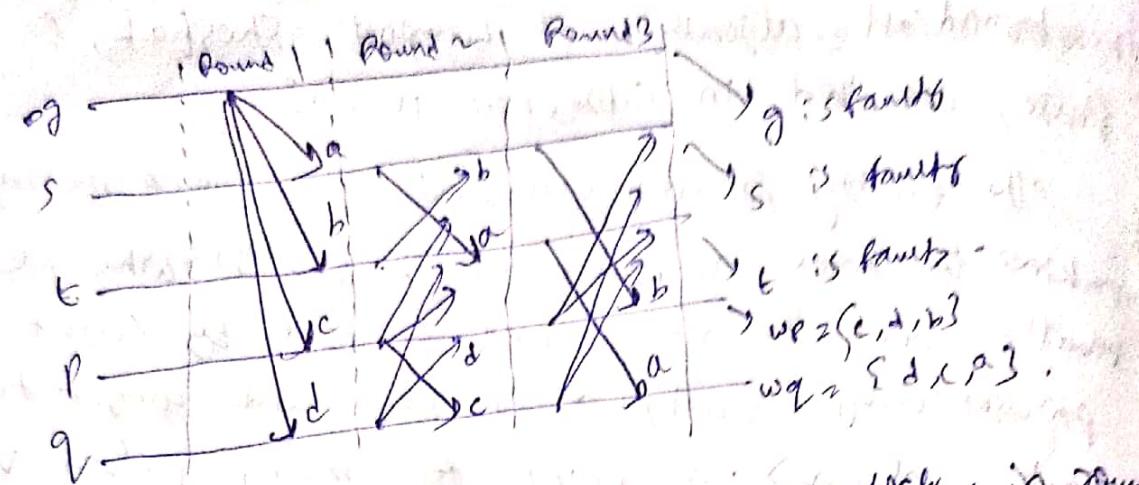
Q15.3 How many messages are exchanged by the broadcast algorithm of Lamport, Shostak, & Pease, described in subsection 15.2.1?

As the faulty processes can send any no. of messages, making the total no. unbounded; this is why we usually count the no. of messages sent by correct processes only; faulty p_1, p_2, \dots, p_n sends $\text{cons}(\text{val}, v)$; p_1, p_2 for many values v and $\text{cons}(\text{val}, v) : p_1, p_2$ for many values v to correct p to trap p into sending many messages. Therefore, a message that repeats another message's sequence of signature is declared invalid.

Then, a correct general sends at most $n-1$ messages in pulse 1, & for each sequence g, p_1, p_2, \dots, p_i in pulse 1, at most $n-i$ messages are sent in pulse 1 by a process (p_i) . Thus, the no. of messages sent by correct processes is at most $\sum_{i=1}^{n-1} (n-i)$.

Q15.4 Is there an execution of the protocol of Dolev & Stora, described in subsection 15.2.1, where correct processes p & q end with $w_p \neq w_q$?

As let p & q be correct; the 1st & 2nd value alerted by p or q are forwarded, & also alerted by the other. So, we let them both receive a 3rd value from faulty process, which is not received by the other.



- P & Q forward the values c , d , e , respectively, in round 2 & that they forward in round 3 their local received value, namely d , e , c , respectively.
- The 3rd received value (bfa , respectively) is not forwarded in the next round.

Q15.5 show that order preserving remaining in the range 1 through N can be solved when interactive consistency is achieved -

- The processes achieve interactive consistency on the vector of inputs; the name of P will be the rank of v_P in this vector.
- Byzantine process may refuse to participate, but their non-default values are assumed to have the higher ranks in the list. They may also cheat by submitting offer values than their input, for eg, the name of a correct process then this name appears multiply in the list & the correct process decides on the smallest; for which (in the sorted vector) $V[i] = v_P$.
- Because all correct processes decide on the same vector & this decision vector contains the RP of each correct process (by dependence), the decisions of correct processes are different & in the same $1 \dots p^N$.

Q15.6 Assume that P signs 2 messages M_1, M_2 with the Elgamal signature scheme, using the same value of a , show how to find P 's secret key from pre & signed messages.

Ans Process P receives signed messages M_1, M_2 from P . i.e., the triples (M_1, δ_1, s_1) & (M_2, δ_2, s_2) , where $\delta_i = g^{a_i} \pmod{P}$ & $s_i = (M_i - \delta_i) a_i^{-1} \pmod{P-1}$.

\rightarrow Range of a , i.e., $a_1 \neq a_2$ is detected by observing that $a_1 \neq a_2$ in this case. Subtracting s_2 from s_1 eliminates the unknown term $a\delta_i$ as $s_1 - s_2 = (M_1 - M_2) a^{-1} \pmod{P-1}$.
 $i.e., a = (M_1 - M_2) (s_1 - s_2)^{-1} \pmod{P-1}$.

\rightarrow Now, we find $d = (M_1 - a s_1)^{-1} \pmod{P-1}$.
Q15.7 Let n be the product of 2 large prime nos, assume that a black box is given that computes square roots. That is, given a quadratic residue y , the box outputs an x with $x^2 \equiv y \pmod{n}$. Show how the box can be used to factor n .

Ans In rings with zero-divisors, i.e., the axiom $a, b \neq 0 \Rightarrow ab \neq 0$ does not hold, a square is the square of at most two elements. Indeed, $x^2 = z^2$ leads $(x-z)(x+z) = 0$. Hence, by quoted axiom, $(x-z) = 0$ or $(x+z) = 0$. i.e. $x = z$ or $x = -z$.

\rightarrow In particular, if $P = \text{prime}$, \mathbb{Z}_P is free of zero-divisors, & $x^2 \equiv y^2 \pmod{P}$ implies $y \equiv \pm x \pmod{P}$.
 \rightarrow Now, let $n = p \cdot q$ be a composite; the ring \mathbb{Z}_n has zero-divisors because $a \cdot b = 0$ holds if p divides a & p divides b , or vice versa. Each square can be the square of upto 4 nos;

if $y \equiv n$, then also $y \equiv z_1^2$ for
 q, s.t. $z_1 \equiv n \pmod{p}$ & $z_1 \equiv y \pmod{q}$
 q, s.t. $z_2 \equiv n \pmod{p}$ & $z_2 \equiv -y \pmod{q}$
 q, s.t. $z_3 \equiv n \pmod{p}$ & $z_3 \equiv y \pmod{q}$
 q, s.t. $z_4 \equiv n \pmod{p}$ & $z_4 \equiv -y \pmod{q}$

The two rows form 2 parts of ∂D^2 and ∂D^2 is such that

$f(z_1 + z_2)$ if possessing z_1 & z_2 such that
 $z_1 + z_2$ is useless. However, if we possess
 z_1 & z_2 from different paths, compute $a = z_1 + z_2$
& $b = z_1 - z_2$ & observe that neither a nor b
equals 0, while $a \cdot b = 0$. Hence, each of a & b
is divided by one of p & q , so $\gcd(a, n) \geq$
a non-trivial factor of n which is coprime to n modulo fully.

- The Salmar - root box can be used to produce such a pair. Select a random no. n & use the box to produce a square root $\pm \sqrt{n^2}$.
- Because n was selected randomly, there is a probability of $\frac{1}{2}$ that n may come from different pairs, generating the factors of n .

(the main work in finding a non-trivial pair of no's
generating the factors of a number by modern factoring methods)

(the main work in modern law,
spent in finding a non-torial part of no's
with same sa.)

Q. 15.8 Show that if clocks C_P & C_Q have δ -bounded drifts and are δ -synchronized at real time t , then they are $\delta(1+\rho)$ -synchronized at clock time $\rightarrow C_P(t)$.

\Rightarrow The clock C_Q has ρ -bounded drifts, meaning that its advance in the real time from t_1 to t_2 (where $t_2 \geq t_1$) differs at most a factor of $1+\rho$ from $t_2 - t_1$, i.e.

$$(t_2 - t_1)(1+\rho) \leq C_Q(t_2) - C_Q(t_1) \leq (t_2 - t_1)(1+\rho)$$

\Rightarrow This implies that the amount of real time in which the clock advances from T_1 to T_2 (where $T_2 \geq T_1$) differs at most a factor of $1+\rho$ from $T_2 - T_1$, i.e.

$$(T_2 - T_1)(1+\rho) \leq C_Q(T_2) - C_Q(T_1) \leq (T_2 - T_1)(1+\rho)$$

\Rightarrow Using only the inequality we find, for any T, T' :

$$|C_Q(T) - C_Q(T')| \leq (1+\rho) |T - T'|$$

\Rightarrow Consequently, for $T = C_P(t)$, we have, $|C_Q(C_P(t)) - C_Q(C_P(t'))| \leq (1+\rho) |C_P(t) - C_P(t')|$ (by the bounded drifts of C_Q as C_P, C_Q are δ -sync. at real time t)

$$\leq (1+\rho) \cdot S$$

159 → the probabilistic clock synchronization protocol is efficient if many messages have a delay close to δ_{\max} (i.e. $F(\delta_{\max} + \epsilon)$ moves away from 0 quickly, even for small ϵ).

→ we have a "dual" protocol that works efficiently if many messages have a delay close to δ_{\min} (i.e. if $F(\delta_{\min} - \epsilon)$)

stays sufficiently far from 1 even for small ϵ).

→ shows that the expected message complexity is bounded by $2 \cdot \frac{1}{1 - F(\delta_{\max} - \epsilon)}$ & more ~~is~~ a bound on the expected running time

→ problem P asks the requestor for the time & the sender

replies immediately with a $(time, t)$ message, & P measures the time T when sending the request & delivering the reply. As the delay of the $(time, t)$

message was at least $\delta - \delta_{\max}$ (this is because at most δ_{\max} of the δ time interval was used by the request message) & at most δ_{\max} , it differs at most $\delta_{\max} - \frac{\delta}{12}$ from T . So, P sets t_{req}

at most $\delta_{\max} - \frac{\delta}{12}$ from $T + \frac{\delta}{12}$, achieving a $\delta_{\max} - \frac{\delta}{12}$ precision; if ϵ is smaller than $\delta_{\max} - \frac{\delta}{12}$, the request is repeated.

→ the probability that the delay of the request (or response) exceeds $\delta_{\max} - \epsilon$ is $1 - F(\delta_{\max} - \epsilon)$,

so the probability that both delays exceed $\delta_{\max} - \epsilon$ is $[1 - F(\delta_{\max} - \epsilon)]^2$. This implies that the expected no. of trials is at most $[1 - F(\delta_{\max} - \epsilon)]^{-2}$.

→ hence, the expected no. of messages is at most $2 \cdot [1 - F(\delta_{\max} - \epsilon)]^{-2}$ & the expected time is at most $2\delta_{\max}[1 - F(\delta_{\max} - \epsilon)]$.

8/15/12 Does the set $A\delta$ in Mgs 15,6 satisfy width
 $(A\delta) \leq 8$ for all correct ρ ?
assume $\rho = \epsilon$ precompute

AP $\leq \delta$ to an order of $n-t$ problems
No, certainly not. Assume $n-t$ problems
Submit the values $n-s$, & in addition I received
the values $n-s$ & $n+\delta$. Although at least one
of these is erroneous (they differ by more than
 δ), neither of them is rejected because both have
 $n-t$ submissions that differs δ from it. This lower bounds
width (AP) to 2δ .

Chapter 17

Q 17.1 Prove:-

- assume that either
- (1) all terminal configurations belong to L ,
 - (2) \exists a function $f: L \rightarrow W$, where W is a well-founded set, & for each partition

$$Y \rightarrow S, f(Y) > f(S) \text{ or } S \in L \text{ holds}$$

then $S = (C, \rightarrow)$ satisfies consistency

as consider an execution $E = (\sigma_0, \sigma_1, \dots)$ of S .

\Rightarrow if E is finite, its last configuration is terminal & belongs to L by (1). otherwise, consider the sequence $(f(\sigma_0), f(\sigma_1), \dots)$; as there are no infinite decreasing sequences in W , there is an i such that $f(\sigma_i) > f(\sigma_{i+1})$ doesn't hold. By (2), $\sigma_{i+1} \in L$.

Q 17.2 prove that Dijkstra's token ring reaches a legitimate configuration in $O(N^2)$ steps. shorten the analysis by giving a single norm function, gradually bounded in N , that decreases with every step of the algorithm.

As in the initial state F is at most $\frac{1}{2}N(N-1)$ & every step of process 0 increases F by at most $N-1$ (if it gives privilege to process 1). So the no. of steps by process 0 is bounded by $\frac{1}{2}N(N-1) / O(N^{-1})$ which brings the total no. of steps to at most $\frac{3}{2}N^{1/2}$.

→ A suitable monom function is defined as follows

[DAS⁹⁴ Sol. §.2]. Call configuration δ single-segment if there exists $j \in N$ such that for all

$$i \in S : \delta_i = \delta_0 \text{ and } i > j : \delta_i = \delta_0.$$

Define, $A(\delta)$ is calculated as follows

$$A(\delta) = \begin{cases} 0, & \text{if } \delta \text{ is single-segment} \\ \text{smallest } i \text{ s.t. } \delta_0 + \\ (\dots + \delta_i) \text{ doesn't occur in } \delta, & \text{otherwise} \end{cases}$$

& $W(\delta)$ by $W(\delta) = N \cdot A(\delta) + F(\delta)$. For each step,

$$\delta \rightarrow \delta' \Rightarrow [W(\delta') \geq W(\delta) \vee \delta' \in L] \text{ holds.}$$

Q174 Modify ring-orientation algorithm (17.1) so that it will terminate after establishing its postcondition.

Ans → modify action (2) so that the idle process will receive the token only if p & q are oriented differently; if they are oriented the same way, the sending process becomes idle.

Actions	1a	1b	2
(1a)	$\frac{S}{P}$	$\frac{S}{P}$	$\frac{S}{P} \rightarrow R$
(1b)	$\frac{S}{P}$	$\frac{S}{P}$	$\frac{S}{P} \rightarrow S$

→ Action (1a) is the same first step as before, but action (2) will kill the token where originally it would be forwarded to a process that already has the same orientation. This ensures that token circulation ceases, yet we are still guaranteed to have tokens as long as the ring is not legitimate.

Q17.6 Show that the maximal-matching algorithm
can be implemented uniformly in the link-registers
read-all model.

Ans → Replace the variable p_{ij} by a 3-valued variable
 c_{pq} for each neighbour q , with value you if p has
selected q , other if p has selected $\neq q$, and none if
 p has not made a selection. These variables satisfy
a local consistency:

$$dc(p) \equiv (\forall q \in \text{Neigh} : c_{pq} = \text{none})$$

$$\wedge (\exists q \in \text{Neigh} : c_{pq} = \text{you}) \wedge \forall q \in \text{Neigh} : c_{pq} = \text{other}).$$

→ An additional "clearing" action c_p , executed at most
once for every process, establishes local consistency,
qualifications below own over Neigh_p .

$$\text{val } c_p : \{(\text{you}, \text{other}, \text{none})\};$$

$$a : \{ -dc(p) \}$$

forall q do $c_{pq} := \text{none}$

$$mp : \{ dc(p) \wedge c_{pq} = \text{none} \wedge (c_{pq} \neq \text{you}) \}$$

forall q do $c_{pq} = \text{other} ; c_{pq} \neq \text{you}$

$$sp : \{ dc(p) \wedge c_{pq} \neq \text{none} \wedge (\forall q \neq p : c_{pq} = \text{none}) \}$$

forall q do $c_{pq} = \text{other} ; c_{pq} = \text{you}$

$$ep : \{ dc(p) \wedge c_{pq} = \text{you} \wedge (c_{pq} \neq \text{other}) \}$$

forall q do $c_{pq} := \text{none}$

→ Each process is locally consistent after almost one step, & the algorithm then behaves like the original one.

Q12.7 Show that the maximal matching algorithm stabilizes in $N^2 + O(N)$ steps & show that there is an $\Omega(N^2)$ lower bound on the (cost - cost)
no. of steps.

\rightarrow From each matching step (action M_p) the value
 of $f + w$ decreases by at least 2, & the other steps
 decrease the value of the second component of F .
 \rightarrow Hence, the function $F_2((E+f+w)/2), 2c+f$ is
 also a room, & provides stabilization within $N^2 \cdot 2^N$

To show the lower bound we inductively construct
 an execution of $\frac{1}{2}k^2$ steps on \mathbb{N} a closure containing
 an even number of processes p_1, p_2, \dots, p_{k-1} , first process p_1 ,
 k (even) free processes p_1, p_2, \dots, p_{k-1} , then process p_k matched p_{k-1} ,
 through p_{k-1} select p_k , then process p_k matched p_{k-2} ,
 through p_{k-2} unchain p_k , & finally process p_k , through p_{k-2} unchain
 & finally processes p_1, p_2, \dots, p_{k-2} , the NW
 from p_k . After these $\frac{1}{2}k^2$ steps, the NW contains
 $k-2$ free processes p_1, p_2, \dots, p_{k-2} , the NW contains
 from p_k . After these $\frac{1}{2}k^2$ steps, the NW contains
 $k-2$ free processes, so by induction the
 $k-2$ free processes, so by induction the
 $k-2$ free processes, so by induction the

one cut on 1st day of first week
Sep 1

Q1A.8 Design a stabilizing algorithm to construct a maximal independent set & compute the maximal no. of steps before stabilization -

Ans Assume the state dead - we model it have a boolean mif for each P_j indicating if P_j is currently in the set M , there are 2 actions :-

(1) if both P_j & a neighbour are in M ,

$mif := \text{false}$;

(2) if neither P_j nor any neighbours is in M ,

$mif := \text{true}$.

\rightarrow Call P_j black if :- (1) applies, gray if :- (2) applies & white if neither applies. A max. is off.

\rightarrow Action (1) makes the black process white without making any other process black (a white neighbour could become gray) & action (2) makes a gray process white without darkening any other node.

\rightarrow So, the pair formed by the no. of black & gray processes decreases lexicographically with every step & after less than N^2 steps a terminal configuration is reached. In a terminal configuration, no 2 neighbours are in M , & every node not in M has a neighbour in M , hence M is a maximal independent set.

Q1A.10 prove that outerplanar graphs are 3-colorable, design a stabilizing algorithm that computes a 3-coloring of an outerplanar graph.

Ans An outerplanar graph has a node of degree 2 or less; so, reducing '6' by '3' & '5' by '2', we get the answer.

Q 17.12 Show that the update algorithm can be used for selection & computation of a ~~BFS~~ spanning tree by giving an appropriate path cost function.

Ans Let the cost of $\pi = (p_0, p_1, \dots, p_N)$ be the path (p_0, k) . Ordering is lexicographically, but to get length increase, paths of length N or longer are larger than paths of length below N .

Q 17.13 Give a stabilizing algorithm for computing the network size.

Ans When a spanning tree is known, the size is computed by summation per subtree i.e. each node assigns its count the sum of the counts of its children plus one. Construction of a spanning tree without a prior knowledge of the network size was solved.

Q 17.14 Show how to compute the depth of a tree with the update algorithm.

Ans The update algorithm can compute the longest upward paths for each node with the following choice of c, f , & the order. Let $c(p)$, for each process p , & let $f(q, n) = n+1$ if p is the father of q & - p otherwise; the order is reversed, i.e., longer paths are smaller than shorter ones.